

# EDA No. 4 AAA Project Martin George mgeorgevienna@gmail.com

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
df = pd.read_csv('member_sample.csv', index_col = 0)
```

## Application of classification model on AAA data

```
In [2]: df.head()
df.info()
df.columns

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21344 entries, 0 to 99998
Columns: 112 entries, Individual Key to Was Towed To AAR Referral
dtypes: float64(35), object(77)
memory usage: 18.4+ MB
```

```
Out[2]: Index(['Individual Key', 'Household Key', 'Member Flag', 'City',
              'State - Grouped', 'ZIP5', 'ZIP9', 'FSV CMSI Flag',
              'FSV Credit Card Flag', 'FSV Deposit Program Flag',
              ...
              'SC Vehicle Manufacturer Name', 'SC Vehicle Model Name',
              'SVC Facility Name', 'SVC Facility Type', 'Total Cost',
              'Tow Destination Latitude', 'Tow Destination Longitude',
              'Tow Destination Name', 'Was Duplicated', 'Was Towed To AAR Referral'],
              dtype='object', length=112)
```

In [3]:

df.head()

Out[3]:

	Individual Key	Household Key	Member Flag	City	State - Grouped	ZIP5	ZIP9	FSV CMSI Flag	FSV Credit Card Flag	FSV Deposit Program Flag	...	SC Vehicle Manufacturer Name	SC Vehicle Model Name	SVC Facility Name	F
0	10000003.0	10462590.0	Y	NEW HAVEN	CT	6511.0	65111349.0	N	N	N	...	NaN	NaN	NaN	
1	52211550.0	4500791.0	Y	WEST WARWICK	RI	2893.0	28933850.0	N	Y	N	...	TOYOTA	CAMRY	ASTRO WRECKER SERVICE	independe
2	52211550.0	4500791.0	Y	WEST WARWICK	RI	2893.0	28933850.0	N	Y	N	...	TOYOTA	CAMRY	Astro Wrecker Service	independe
3	52211550.0	4500791.0	Y	WEST WARWICK	RI	2893.0	28933850.0	N	Y	N	...	TOYOTA	CAMRY	ASTRO WRECKER SERVICE	independe
4	52211550.0	4500791.0	Y	WEST WARWICK	RI	2893.0	28933850.0	N	Y	N	...	TOYOTA	CAMRY	ASTRO WRECKER SERVICE	independe

5 rows × 112 columns

In [4]:

df['Do Not Direct Mail Solicit'].isna().sum()

Out[4]:

1

In [5]:

df.groupby('Do Not Direct Mail Solicit')['Total Cost'].mean()

Out[5]:

Do Not Direct Mail Solicit  
0.0 47.270066  
1.0 47.363815  
Name: Total Cost, dtype: float64

```
In [6]: df[['Do Not Direct Mail Solicit', 'Email Available', 'Email Status']].head()
```

```
Out[6]:
```

	Do Not Direct Mail Solicit	Email Available	Email Status
0	NaN	NaN	NaN
1	0.0	0.0	NaN
2	0.0	0.0	NaN
3	0.0	0.0	NaN
4	0.0	0.0	NaN

```
In [7]: sub_df = df[df['Email Available'] == 1.0]
```

```
In [8]: sub_df.shape
```

```
Out[8]: (11442, 112)
```

```
In [9]: sub_df['Email Status'].value_counts(dropna = False)
```

```
Out[9]: Active          4394
Unsubscribed          4376
NaN                   2031
Held                   515
Bounced               126
Name: Email Status, dtype: int64
```

```
In [10]: #df.loc[(df['Email Available'] != 1.0 )]
df_email_miss= df.loc[(df['Do Not Direct Mail Solicit'] != 1.0 ) & (df['Email Available'] != 1.0 ) ]

u_hh = df_email_miss["Household Key"].unique()
len(u_hh)
#n_by_state = df.groupby("state")["Last_name"].count()
```

```
Out[10]: 3836
```

```
In [11]: df_mail=df[df['Do Not Direct Mail Solicit'] != 1.0]
df_mail_ava=df_mail[df['Email Available'] != 1.0]
df_mail_ava.shape
```

C:\Users\george\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel\_launcher.py:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
Out[11]: (9221, 112)
```

```
In [12]: df_mail_ava.shape
```

```
Out[12]: (9221, 112)
```

## Recommendation for AAA on email address

There are 9221 rows with email addresses are missing , even though these c ustomers have not blocked direct emails from AAA. Out of these 9221, 3836 are unique householders

```
In [13]: df['Reason Joined'].shape
```

```
Out[13]: (21344,)
```

```
In [14]: df['Reason Joined'].value_counts(dropna = False)
```

```
Out[14]: NaN                20956  
U                  168  
Dependable Services    127  
5                     45  
Family Plan Avail      19  
Nation Wide Rd Srv      7  
Gift Membership         5  
Free Membership         4  
Club Reputation         3  
3                      3  
Convenient Offices      1  
Variety of Services     1  
Direct Mail             1  
Prior Family Exp        1  
7                      1  
Other                   1  
Recommend/Referral      1  
Name: Reason Joined, dtype: int64
```

## Study on 'Reason Joined' field

Using KNN algorithm, we can fill the missing values based on another field, example 'Total Cost'

```
In [15]: test_df = df.loc[df['Reason Joined'].isna()]
```

```
In [16]: test_df.shape
```

```
Out[16]: (20956, 112)
```

```
In [17]: train_df = df.loc[df['Reason Joined'].notna()]
```

```
In [18]: train_df.shape
```

```
Out[18]: (388, 112)
```

```
In [19]: train_df['Reason Joined'].value_counts(dropna = False)
```

```
Out[19]: U                168
Dependable Services    127
5                      45
Family Plan Avail      19
Nation Wide Rd Srv      7
Gift Membership         5
Free Membership         4
3                       3
Club Reputation         3
Other                   1
Convenient Offices      1
Recommend/Referral      1
7                       1
Variety of Services     1
Direct Mail             1
Prior Family Exp        1
Name: Reason Joined, dtype: int64
```

```
In [20]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [21]: knn = KNeighborsClassifier()
```

```
In [22]: train_df = train_df[['Total Cost', 'Reason Joined']].dropna()
```

```
In [23]: X= train_df[['Total Cost']]
y = train_df['Reason Joined']
```

```
In [24]: X.shape
```

```
Out[24]: (277, 1)
```

```
In [25]: y.shape
```

```
Out[25]: (277,)
```

```
In [26]: knn.fit(X,y)
```

```
Out[26]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                             weights='uniform')
```

```
In [27]: fill_x = test_df[['Total Cost', 'Reason Joined']].dropna(subset = ['Total Cost'])
```

```
In [28]: fill_x['Reason Joined'].value_counts(dropna = False)
```

```
Out[28]: NaN      13667  
         Name: Reason Joined, dtype: int64
```

```
In [29]: fill_x.shape
```

```
Out[29]: (13667, 2)
```

```
In [30]: test_df.shape
```

```
Out[30]: (20956, 112)
```

```
In [31]: x_test = fill_x[['Total Cost']]
```

```
In [32]: knn.predict(x_test)
```

```
Out[32]: array(['Dependable Services', 'Dependable Services',  
                'Dependable Services', ..., 'Dependable Services',  
                'Dependable Services', 'Dependable Services'], dtype=object)
```

```
In [33]: x_test.shape
```

```
Out[33]: (13667, 1)
```

```
In [34]: fill_b = fill_x
```

```
In [35]: fill_x['Reason Joined'].value_counts(dropna = False)
```

```
Out[35]: NaN      13667  
         Name: Reason Joined, dtype: int64
```

```
In [36]: fill_x['Reason Joined'] = knn.predict(x_test)
```

```
In [37]: fill_x[['Total Cost', 'Reason Joined']]
```



Out[37]:

	Total Cost	Reason Joined
1	32.50	Dependable Services
2	30.00	Dependable Services
3	32.50	Dependable Services
4	30.00	Dependable Services
5	53.00	Dependable Services
6	30.00	Dependable Services
7	32.00	Dependable Services
8	32.00	Dependable Services
9	32.50	Dependable Services
11	58.85	Dependable Services
12	53.00	Dependable Services
13	53.00	Dependable Services
16	53.00	Dependable Services
17	53.00	Dependable Services
18	53.00	Dependable Services
19	29.00	Dependable Services
20	28.00	U
21	53.00	Dependable Services
22	53.00	Dependable Services
23	53.00	Dependable Services
24	53.00	Dependable Services
25	53.00	Dependable Services
26	58.85	Dependable Services
27	53.00	Dependable Services
28	53.00	Dependable Services
29	53.00	Dependable Services
32	58.85	Dependable Services
34	58.85	Dependable Services
36	29.00	Dependable Services

	Total Cost	Reason Joined
41	33.00	U
...	...	...
99515	53.00	Dependable Services
99516	53.00	Dependable Services
99539	39.00	U
99540	58.85	Dependable Services
99541	0.00	U
99542	53.00	Dependable Services
99543	53.00	Dependable Services
99544	53.00	Dependable Services
99545	53.00	Dependable Services
99966	58.85	Dependable Services
99967	53.00	Dependable Services
99969	47.00	Dependable Services
99970	58.85	Dependable Services
99971	58.85	Dependable Services
99972	53.00	Dependable Services
99973	53.00	Dependable Services
99974	58.85	Dependable Services
99977	34.00	U
99983	29.00	Dependable Services
99984	44.00	Dependable Services
99985	29.00	Dependable Services
99986	47.00	Dependable Services
99987	44.00	Dependable Services
99988	82.00	Dependable Services
99989	29.43	Dependable Services
99993	53.00	Dependable Services
99994	36.00	U
99995	53.00	Dependable Services

	Total Cost	Reason Joined
99996	58.85	Dependable Services
99997	58.85	Dependable Services

13667 rows × 2 columns

```
In [38]: fill_x['Reason Joined'].value_counts(dropna = False)
```

```
Out[38]: Dependable Services    10679
U                                2840
5                                 144
Family Plan Avail                4
Name: Reason Joined, dtype: int64
```

```
In [39]: fill_b['Reason Joined'].value_counts(dropna = False)
```

```
Out[39]: Dependable Services    10679
U                                2840
5                                 144
Family Plan Avail                4
Name: Reason Joined, dtype: int64
```

**KNN is an optional algorithm to fill the missing values of a feature column, based on an another reliable feature. In this aabove example, based on the Total cost, we could predict the missing values of " Reason Joined" feature.**

**sklearn imputer is also a method for filling missing values. Testing follows.**

```
In [137]: from sklearn.impute import SimpleImputer
```

```
In [149]: imputer = SimpleImputer(strategy = 'median')
```

```
In [146]: df_to_impute = df.select_dtypes(['int', 'float'])
```

```
In [147]: df_to_impute
```

Out[147]:

	Individual Key	Household Key	ZIP5	ZIP9	Length Of Residence	Do Not Direct Mail Solicit	Email Available	ERS ENT Count Year 1	ERS ENT Count Year 2	ERS ENT Count Year 3	...	Member Match Flag	Member Number and Associate ID	Plus Cost	Premier Cost
0	10000003.0	10462590.0	6511.0	65111349.0	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
1	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
2	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
3	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
4	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
5	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
6	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
7	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
8	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
9	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	15300.0	0.0	0.0
10	1606764.0	4317516.0	2878.0	28781026.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
11	2766867.0	11622991.0	2889.0	28892920.0	13.0	0.0	1.0	1.0	2.0	0.0	...	1.0	16300.0	0.0	0.0
12	2766867.0	11622991.0	2889.0	28892920.0	13.0	0.0	1.0	1.0	2.0	0.0	...	1.0	16300.0	0.0	0.0
13	2766867.0	11622991.0	2889.0	28892920.0	13.0	0.0	1.0	1.0	2.0	0.0	...	1.0	16300.0	0.0	0.0
14	2766869.0	11622991.0	2889.0	28892920.0	13.0	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
15	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN
16	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	1.0	16322.0	0.0	0.0
17	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	1.0	16322.0	0.0	0.0
18	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	1.0	16322.0	0.0	0.0
19	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	1.0	16322.0	0.0	0.0
20	13746947.0	579810.0	2863.0	28631322.0	15.0	0.0	1.0	0.0	0.0	0.0	...	1.0	18200.0	0.0	0.0
21	1788453.0	7187017.0	2888.0	28882811.0	15.0	0.0	1.0	0.0	0.0	0.0	...	1.0	18800.0	0.0	0.0
22	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	1.0	18820.0	0.0	0.0
23	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	1.0	18820.0	0.0	0.0
24	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	1.0	18820.0	0.0	0.0
25	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	1.0	18820.0	0.0	0.0
26	1788455.0	7187017.0	2888.0	28882811.0	15.0	0.0	1.0	1.0	0.0	0.0	...	1.0	18821.0	0.0	0.0

	Individual Key	Household Key	ZIP5	ZIP9	Length Of Residence	Do Not Direct Mail Solicit	Email Available	ERS ENT Count Year 1	ERS ENT Count Year 2	ERS ENT Count Year 3	...	Member Match Flag	Member Number and Associate ID	Plus Cost	Premier Cost	
27	14243585.0	7728088.0	2806.0	28065003.0	15.0	0.0	1.0	0.0	0.0	1.0	...	1.0	19100.0	0.0	0.0	
28	14243587.0	7728088.0	2806.0	28065003.0	13.0	0.0	1.0	0.0	0.0	0.0	...	1.0	19120.0	0.0	0.0	3
29	14243587.0	7728088.0	2806.0	28065003.0	13.0	0.0	1.0	0.0	0.0	0.0	...	1.0	19120.0	0.0	0.0	4
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
99968	4458026.0	1588987.0	2886.0	28861711.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99969	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	1.0	54348200.0	0.0	0.0	20
99970	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	1.0	54348200.0	0.0	0.0	19
99971	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	1.0	54348200.0	0.0	0.0	19
99972	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	1.0	54348200.0	0.0	0.0	4
99973	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	1.0	54348200.0	0.0	0.0	7
99974	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	1.0	54348200.0	0.0	0.0	16
99975	12849941.0	16604128.0	2906.0	29063709.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99976	22426406.0	45466286.0	2809.0	28092304.0	12.0	0.0	1.0	0.0	0.0	0.0	...	NaN	NaN	NaN	NaN	
99977	22426405.0	45466286.0	2809.0	28092304.0	15.0	0.0	0.0	0.0	1.0	0.0	...	1.0	54349620.0	0.0	0.0	4
99979	19764804.0	15397653.0	2871.0	28712143.0	NaN	0.0	1.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99980	19764802.0	15397653.0	2871.0	28712143.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99981	19764801.0	15397653.0	2871.0	28712143.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99982	19764793.0	15397653.0	2871.0	28712143.0	0.0	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99983	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	1.0	54353700.0	0.0	0.0	6
99984	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	1.0	54353700.0	18.0	0.0	6
99985	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	1.0	54353700.0	0.0	0.0	7
99986	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	1.0	54353700.0	18.0	0.0	6
99987	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	1.0	54353700.0	18.0	0.0	4
99988	16521336.0	13735475.0	2809.0	28091350.0	10.0	0.0	0.0	0.0	0.0	0.0	...	1.0	54353720.0	39.0	0.0	4
99989	1619870.0	5462399.0	2879.0	28791421.0	0.0	0.0	1.0	0.0	0.0	0.0	...	1.0	54365300.0	0.0	0.0	19
99990	1619868.0	5462399.0	2879.0	28791421.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99991	1619869.0	5462399.0	2879.0	28791421.0	NaN	0.0	0.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	
99992	54745437.0	5462399.0	2816.0	28167132.0	10.0	0.0	0.0	0.0	0.0	0.0	...	NaN	NaN	NaN	NaN	

	Individual Key	Household Key	ZIP5	ZIP9	Length Of Residence	Do Not Direct Mail Solicit	Email Available	ERS ENT Count Year 1	ERS ENT Count Year 2	ERS ENT Count Year 3	...	Member Match Flag	Member Number and Associate ID	Plus Cost	Premier Cost	
<b>99993</b>	25797262.0	20330346.0	2886.0	28867552.0	NaN	0.0	0.0	1.0	0.0	3.0	...	1.0	54367300.0	0.0	0.0	4
<b>99994</b>	25797262.0	20330346.0	2886.0	28867552.0	NaN	0.0	0.0	1.0	0.0	3.0	...	1.0	54367300.0	6.0	0.0	4
<b>99995</b>	25797262.0	20330346.0	2886.0	28867552.0	NaN	0.0	0.0	1.0	0.0	3.0	...	1.0	54367300.0	0.0	0.0	4
<b>99996</b>	28273400.0	8325571.0	2886.0	28868235.0	13.0	0.0	1.0	2.0	0.0	3.0	...	1.0	54369500.0	0.0	0.0	19
<b>99997</b>	28273400.0	8325571.0	2886.0	28868235.0	13.0	0.0	1.0	2.0	0.0	3.0	...	1.0	54369500.0	0.0	0.0	19
<b>99998</b>	28273400.0	8325571.0	2886.0	28868235.0	13.0	0.0	1.0	2.0	0.0	3.0	...	1.0	54369500.0	NaN	NaN	19

21344 rows × 35 columns



```
In [161]: X = df_to_impute.dropna(subset = ['Total Cost']).drop('Total Cost', axis = 1)
```

```
In [177]: y = df_to_impute.dropna(subset = ['Total Cost'])['Total Cost']
```

```
In [178]: imputer.fit(X)
```

```
Out[178]: SimpleImputer(copy=True, fill_value=None, missing_values=nan,
strategy='median', verbose=0)
```

```
In [179]: X = imputer.transform(X)
```

In [180]: `pd.DataFrame(X)`




Out[180]:

	0	1	2	3	4	5	6	7	8	9	...	23	24	25	26	27	28	29	
0	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	9707320.0	41.00000	-71.0000
1	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	6361198.0	0.00000	0.0000
2	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	9127495.0	0.00000	0.0000
3	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	1593215.0	0.00000	0.0000
4	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	3652711.0	0.00000	0.0000
5	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	1.0	1.0	15300.0	0.0	0.0	4646305.0	0.00000	0.0000
6	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	7373094.0	0.00000	0.0000
7	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	7460803.0	0.00000	0.0000
8	52211550.0	4500791.0	2893.0	28933850.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	15300.0	0.0	0.0	9152542.0	0.00000	0.0000
9	2766867.0	11622991.0	2889.0	28892920.0	13.0	0.0	1.0	1.0	2.0	0.0	...	0.0	1.0	16300.0	0.0	0.0	18842978.0	0.00000	0.0000
10	2766867.0	11622991.0	2889.0	28892920.0	13.0	0.0	1.0	1.0	2.0	0.0	...	0.0	1.0	16300.0	0.0	0.0	2075597.0	0.00000	0.0000
11	2766867.0	11622991.0	2889.0	28892920.0	13.0	0.0	1.0	1.0	2.0	0.0	...	1.0	1.0	16300.0	0.0	0.0	5005248.0	0.00000	0.0000
12	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	16322.0	0.0	0.0	1640562.0	0.00000	0.0000
13	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	16322.0	0.0	0.0	1711324.0	0.00000	0.0000
14	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	16322.0	0.0	0.0	1852044.0	0.00000	0.0000
15	2766868.0	11622991.0	2889.0	28892920.0	5.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	16322.0	0.0	0.0	5441060.0	0.00000	0.0000
16	13746947.0	579810.0	2863.0	28631322.0	15.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	18200.0	0.0	0.0	163014.0	0.00000	0.0000
17	1788453.0	7187017.0	2888.0	28882811.0	15.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	18800.0	0.0	0.0	7711905.0	0.00000	0.0000
18	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	0.0	1.0	18820.0	0.0	0.0	1575870.0	0.00000	0.0000
19	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	0.0	1.0	18820.0	0.0	0.0	1677303.0	0.00000	0.0000
20	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	0.0	1.0	18820.0	0.0	0.0	3947898.0	0.00000	0.0000
21	1788452.0	7187017.0	2888.0	28882811.0	15.0	0.0	0.0	0.0	0.0	0.0	...	0.0	1.0	18820.0	0.0	0.0	3952912.0	0.00000	0.0000
22	1788455.0	7187017.0	2888.0	28882811.0	15.0	0.0	1.0	1.0	0.0	0.0	...	0.0	1.0	18821.0	0.0	0.0	19656374.0	41.71028	-71.4910
23	14243585.0	7728088.0	2806.0	28065003.0	15.0	0.0	1.0	0.0	0.0	1.0	...	1.0	1.0	19100.0	0.0	0.0	385595.0	0.00000	0.0000
24	14243587.0	7728088.0	2806.0	28065003.0	13.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	19120.0	0.0	0.0	3121195.0	41.00000	-71.0000
25	14243587.0	7728088.0	2806.0	28065003.0	13.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	19120.0	0.0	0.0	4997740.0	0.00000	0.0000
26	4064211.0	8724100.0	2879.0	28791619.0	15.0	0.0	1.0	0.0	1.0	0.0	...	0.0	1.0	20500.0	0.0	0.0	17543023.0	0.00000	0.0000
27	195833722.0	12394451.0	2920.0	29205904.0	15.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	20800.0	0.0	0.0	20438910.0	0.00000	0.0000
28	7182300.0	300071.0	6111.0	61111421.0	15.0	0.0	0.0	1.0	0.0	0.0	...	0.0	1.0	20900.0	0.0	0.0	1892106.0	0.00000	0.0000

	0	1	2	3	4	5	6	7	8	9	...	23	24	25	26	27	28	29	
29	53136272.0	37549482.0	2814.0	28140051.0	10.0	0.0	1.0	0.0	1.0	0.0	...	0.0	1.0	28200.0	0.0	0.0	17183027.0	41.92437	-71.674
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
13914	205158507.0	95100521.0	2871.0	28714086.0	15.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	53449000.0	0.0	0.0	4583123.0	0.00000	0.000
13915	205158507.0	95100521.0	2871.0	28714086.0	15.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	53449000.0	0.0	0.0	1777244.0	0.00000	0.000
13916	56389710.0	4641147.0	2886.0	28863516.0	15.0	0.0	1.0	0.0	2.0	3.0	...	0.0	1.0	53485600.0	6.0	0.0	18794375.0	41.74966	-71.423
13917	56389710.0	4641147.0	2886.0	28863516.0	15.0	0.0	1.0	0.0	2.0	3.0	...	0.0	1.0	53485600.0	0.0	0.0	18670374.0	0.00000	0.000
13918	56389710.0	4641147.0	2886.0	28863516.0	15.0	0.0	1.0	0.0	2.0	3.0	...	0.0	1.0	53485600.0	0.0	0.0	5357333.0	0.00000	0.000
13919	8867560.0	4641147.0	2886.0	28863516.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	53485620.0	0.0	0.0	2910805.0	0.00000	0.000
13920	8867560.0	4641147.0	2886.0	28863516.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	53485620.0	0.0	0.0	2328658.0	0.00000	0.000
13921	8867560.0	4641147.0	2886.0	28863516.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	53485620.0	0.0	0.0	5263118.0	0.00000	0.000
13922	8867560.0	4641147.0	2886.0	28863516.0	15.0	0.0	0.0	0.0	0.0	2.0	...	0.0	1.0	53485620.0	0.0	0.0	6173883.0	0.00000	0.000
13923	4458025.0	1588987.0	2886.0	28861711.0	15.0	0.0	0.0	1.0	1.0	0.0	...	0.0	1.0	54333400.0	0.0	0.0	17851393.0	0.00000	0.000
13924	4458025.0	1588987.0	2886.0	28861711.0	15.0	0.0	0.0	1.0	1.0	0.0	...	0.0	1.0	54333400.0	0.0	0.0	3952261.0	41.00000	-71.000
13925	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	0.0	1.0	54348200.0	0.0	0.0	20218515.0	41.75776	-71.404
13926	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	0.0	1.0	54348200.0	0.0	0.0	19306815.0	0.00000	0.000
13927	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	0.0	1.0	54348200.0	0.0	0.0	19164491.0	0.00000	0.000
13928	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	0.0	1.0	54348200.0	0.0	0.0	5280579.0	0.00000	0.000
13929	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	0.0	1.0	54348200.0	0.0	0.0	7338895.0	0.00000	0.000
13930	12849942.0	16604128.0	2906.0	29063709.0	15.0	0.0	0.0	3.0	2.0	0.0	...	0.0	1.0	54348200.0	0.0	0.0	16625008.0	0.00000	0.000
13931	22426405.0	45466286.0	2809.0	28092304.0	15.0	0.0	0.0	0.0	1.0	0.0	...	0.0	1.0	54349620.0	0.0	0.0	5880022.0	0.00000	0.000
13932	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	54353700.0	0.0	0.0	6416688.0	41.00000	-71.000
13933	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	54353700.0	18.0	0.0	6307512.0	41.00000	-71.000
13934	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	54353700.0	0.0	0.0	7027924.0	41.00000	-71.000
13935	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	54353700.0	18.0	0.0	6837393.0	41.00000	-71.000
13936	16521338.0	13735475.0	2809.0	28091350.0	10.0	0.0	1.0	0.0	0.0	0.0	...	0.0	1.0	54353700.0	18.0	0.0	3042243.0	41.00000	-71.000
13937	16521336.0	13735475.0	2809.0	28091350.0	10.0	0.0	0.0	0.0	0.0	0.0	...	0.0	1.0	54353720.0	39.0	0.0	3017750.0	41.00000	-71.000
13938	1619870.0	5462399.0	2879.0	28791421.0	0.0	0.0	1.0	0.0	0.0	0.0	...	1.0	1.0	54365300.0	0.0	0.0	19887931.0	41.45056	-71.525
13939	25797262.0	20330346.0	2886.0	28867552.0	15.0	0.0	0.0	1.0	0.0	3.0	...	0.0	1.0	54367300.0	0.0	0.0	2284969.0	0.00000	0.000
13940	25797262.0	20330346.0	2886.0	28867552.0	15.0	0.0	0.0	1.0	0.0	3.0	...	0.0	1.0	54367300.0	6.0	0.0	3623739.0	41.00000	-71.000
13941	25797262.0	20330346.0	2886.0	28867552.0	15.0	0.0	0.0	1.0	0.0	3.0	...	0.0	1.0	54367300.0	0.0	0.0	3743487.0	0.00000	0.000

	0	1	2	3	4	5	6	7	8	9	...	23	24	25	26	27	28	29	
<b>13942</b>	28273400.0	8325571.0	2886.0	28868235.0	13.0	0.0	1.0	2.0	0.0	3.0	...	0.0	1.0	54369500.0	0.0	0.0	19404669.0	0.00000	0.000
<b>13943</b>	28273400.0	8325571.0	2886.0	28868235.0	13.0	0.0	1.0	2.0	0.0	3.0	...	0.0	1.0	54369500.0	0.0	0.0	19614013.0	0.00000	0.000

13944 rows × 33 columns



In [181]: `pd.DataFrame(X).isna().sum()`

Out[181]:

```
0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
11     0
12     0
13     0
14     0
15     0
16     0
17     0
18     0
19     0
20     0
21     0
22     0
23     0
24     0
25     0
26     0
27     0
28     0
29     0
30     0
31     0
32     0
```

dtype: int64

```
In [186]: from sklearn.pipeline import Pipeline
          from sklearn.linear_model import LinearRegression
          from sklearn.feature_selection import SelectKBest
```

```
In [188]: pipe = Pipeline([('imputer',SimpleImputer()),
                           ('kbest',SelectKBest()),
                           ('lgr',LinearRegression())])
```

```
In [184]: pipe.fit(X,y)
```

```
Out[184]: Pipeline(memory=None,
                  steps=[('imputer', SimpleImputer(copy=True, fill_value=None, missing_values=nan, strategy='mean',
              verbose=0)), ('lgr', LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
              normalize=False))])
```

**KNN is better than mean or median imputer if more than 50% data is missing**

```
In [190]: df['ZIP5'].value_counts() > 2
```

```
Out[190]: 2920.0    True
          2886.0    True
          2889.0    True
          2888.0    True
          2816.0    True
          2919.0    True
          2818.0    True
          2852.0    True
          2904.0    True
          2910.0    True
          2893.0    True
          2864.0    True
          2906.0    True
          2882.0    True
          2879.0    True
          2911.0    True
          2908.0    True
          2915.0    True
          2865.0    True
          2905.0    True
          2914.0    True
          2871.0    True
          2838.0    True
          2840.0    True
          2860.0    True
          2861.0    True
          2806.0    True
          2828.0    True
          2917.0    True
          2921.0    True
          ...
          6401.0   False
          6417.0   False
          6032.0   False
          6897.0   False
          6516.0   False
          6033.0   False
          6405.0   False
          6902.0   False
          6109.0   False
          6413.0   False
          6614.0   False
          6420.0   False
          6092.0   False
          6770.0   False
          6480.0   False
          2877.0   False
```

6605.0	False
6604.0	False
6825.0	False
6438.0	False
6512.0	False
6066.0	False
6484.0	False
6001.0	False
6517.0	False
6333.0	False
6019.0	False
6374.0	False
6119.0	False
2802.0	False

Name: ZIP5, Length: 182, dtype: bool

```
In [189]: df['ZIP5'].value_counts()[df['ZIP5'].value_counts() > 2]
```



```
Out[189]: 2920.0    1518
          2886.0    1141
          2889.0     826
          2888.0     823
          2816.0     770
          2919.0     703
          2818.0     661
          2852.0     648
          2904.0     641
          2910.0     635
          2893.0     606
          2864.0     560
          2906.0     535
          2882.0     531
          2879.0     507
          2911.0     498
          2908.0     475
          2915.0     440
          2865.0     402
          2905.0     378
          2914.0     371
          2871.0     328
          2838.0     322
          2840.0     321
          2860.0     319
          2861.0     317
          2806.0     313
          2828.0     288
          2917.0     284
          2921.0     273
          ...
          6250.0      4
          6082.0      4
          6840.0      4
          6058.0      4
          2808.0      4
          2801.0      4
          6831.0      4
          6416.0      4
          2862.0      4
          6084.0      4
          6384.0      4
          6877.0      3
          6340.0      3
          6415.0      3
          6108.0      3
          6460.0      3
```

```

6234.0    3
6457.0    3
6419.0    3
6880.0    3
6443.0    3
6807.0    3
6357.0    3
6525.0    3
6798.0    3
6067.0    3
6354.0    3
6243.0    3
6277.0    3
6057.0    3
Name: ZIP5, Length: 133, dtype: int64

```

```
In [191]: df.groupby('FSV CMSI Flag').mean()
```

Out[191]:

	Individual Key	Household Key	ZIP5	ZIP9	Length Of Residence	Do Not Direct Mail Solicit	Email Available	ERS ENT Count Year 1	ERS ENT Count Year 2	ERS ENT Count Year 3	...	Member Match Flag	Me Numbe Associ
FSV CMSI Flag													
N	3.403291e+07	1.600860e+07	2947.671848	2.948020e+07	11.552839	0.054041	0.52604	0.517824	0.921864	0.952447	...	1.0	1.091986
Y	2.398762e+07	1.515128e+07	2885.457413	2.885794e+07	11.088766	0.027340	0.75184	0.531746	1.193878	1.090703	...	1.0	1.071187

2 rows × 35 columns

# Python library SMOTE can be also used to populate samples

```
In [192]: df['FSV CMSI Flag'].value_counts(normalize = True)
```

```

Out[192]: N    0.955444
          Y    0.044556
Name: FSV CMSI Flag, dtype: float64

```

```
In [193]: yes = df.loc[df['FSV CMSI Flag'] == 'Y']
```

```
In [194]: no = df.loc[df['FSV CMSI Flag'] == 'N']
```

```
In [195]: no_sample = no.sample(951)
```

```
In [213]: combined = pd.concat( [yes, no_sample]).dropna(subset = ['Total Cost'])
```

```
In [226]: X = combined[['Total Cost']]  
y = combined[['FSV CMSI Flag']]
```

```
In [227]: from sklearn.linear_model import LogisticRegression
```

```
In [228]: lgr = LogisticRegression()
```

```
In [229]: lgr.fit(X,y)
```

C:\Users\george\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\george\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\utils\validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

y = column\_or\_1d(y, warn=True)

```
Out[229]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
    intercept_scaling=1, max_iter=100, multi_class='warn',  
    n_jobs=None, penalty='l2', random_state=None, solver='warn',  
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [231]: lgr.predict_proba(X)
```

```
Out[231]: array([[0.45883936, 0.54116064],  
    [0.43457182, 0.56542818],  
    [0.44661667, 0.55338333],  
    ...,  
    [0.43900919, 0.56099081],  
    [0.43900919, 0.56099081],  
    [0.43457182, 0.56542818]])
```

```
In [232]: lgr.score(X,y)
```

```
Out[232]: 0.5565280816921955
```

**A prediction algorithm can be evaluated by considering its predict probability and score.**

In [ ]: